

Definition :

- This dataset provides a detailed view of the product catalog and pricing structure of Zepto, a fast-growing 10-minute grocery delivery platform. The data captures essential attributes for over 3,000+ SKUs (Stock Keeping Units) across various categories like Fruits & Vegetables, Dairy, Packaged Foods, Beverages, and more

The data is structured to support various types of retail analysis, including:

- Discount trends by category
- Inventory availability and stock-outs
- Price distribution and pricing strategy
- Product naming patterns (suitable for word cloud or NLP tasks)

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px

from wordcloud import wordcloud
```

```
In [2]: df=pd.read_csv('zepto_v2.csv',encoding="ISO-8859-1")
df.head()
```

```
Out[2]:
```

	Category	name	mrp	discountPercent	availableQuantity	discountedSellingPrice
0	Fruits & Vegetables	Onion	2500	16	3	2100
1	Fruits & Vegetables	Tomato Hybrid	4200	16	3	3500
2	Fruits & Vegetables	Tender Coconut	5100	15	3	4300
3	Fruits & Vegetables	Coriander Leaves	2000	15	3	1700
4	Fruits & Vegetables	Ladies Finger	1400	14	3	1200

Data preprocessing

```
In [3]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3732 entries, 0 to 3731
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Category                             3732 non-null   object
1   name                                 3732 non-null   object
2   mrp                                  3732 non-null   int64
3   discountPercent                      3732 non-null   int64
4   availableQuantity                   3732 non-null   int64
5   discountedSellingPrice               3732 non-null   int64
6   weightInGms                         3732 non-null   int64
7   outOfStock                          3732 non-null   bool
8   quantity                            3732 non-null   int64
dtypes: bool(1), int64(6), object(2)
memory usage: 237.0+ KB

```

```
In [4]: df.shape
```

```
Out[4]: (3732, 9)
```

```
In [5]: df.duplicated().sum()
```

```
Out[5]: 2
```

```
In [6]: df.isna().sum()
```

```

Out[6]: Category          0
        name              0
        mrp               0
        discountPercent    0
        availableQuantity  0
        discountedSellingPrice 0
        weightInGms        0
        outOfStock         0
        quantity           0
        dtype: int64

```

```


In [15]: plt.figure(figsize=(12,10))
        sns.countplot(x='Category',
                      data=df,
                      width=0.4,
                      palette='dark:#5A9_r')
        plt.title('Zepto food category')
        plt.xticks(rotation='vertical');

```

C:\Users\ashif\AppData\Local\Temp\ipykernel_13332\128187932.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='Category',
```

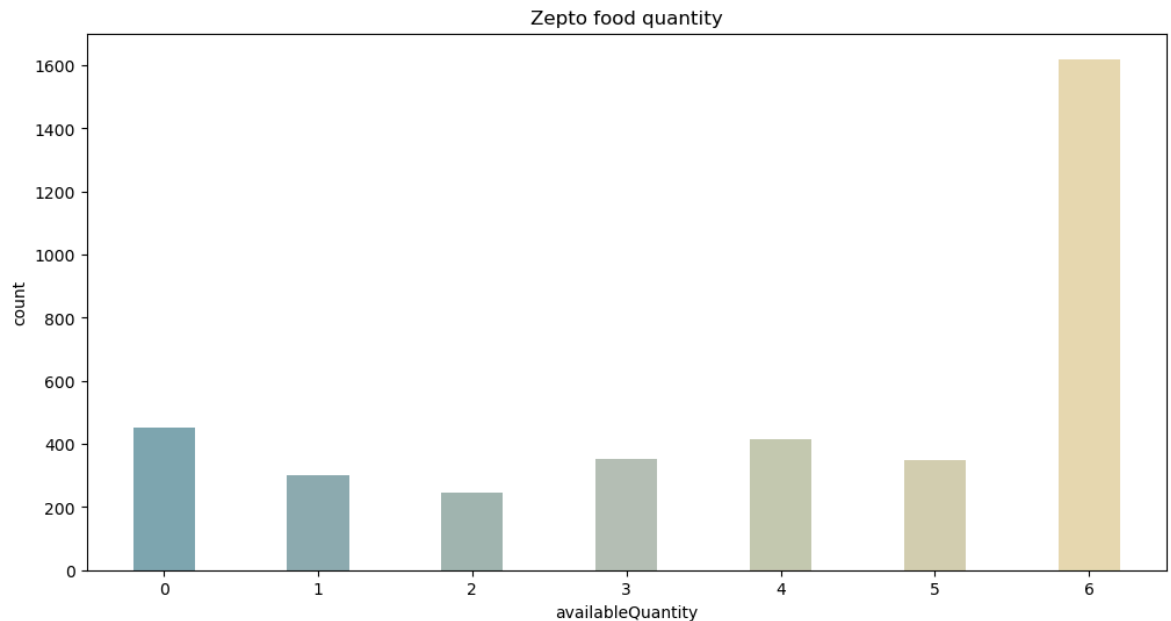


```
In [8]: plt.figure(figsize=(12,6))
sns.countplot(x='availableQuantity',
              data=df,
              width=0.4,
              palette="blend:#7AB,#EDA")
plt.title('Zepto food quantity')
plt.xticks(rotation='horizontal');
```

C:\Users\ashif\AppData\Local\Temp\ipykernel_13332\322488579.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='availableQuantity',
```




```
In [17]: plt.figure(figsize=(8,4))
sns.countplot(x='outOfStock',
              data=df,
              palette='ch:s=.25,rot=-.25',
              width=0.4)
plt.legend()
plt.title('Zepto out of stock');
```

C:\Users\ashif\AppData\Local\Temp\ipykernel_13332\1871124277.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='outOfStock',
C:\Users\ashif\AppData\Local\Temp\ipykernel_13332\1871124277.py:6: UserWarning: No artists with labels found to put in legend. Note that artists whose label start with an underscore are ignored when legend() is called with no argument.
plt.legend()
```



In [10]: `df.head()`

Out[10]:

	Category	name	mrp	discountPercent	availableQuantity	discountedSellingPrice
0	Fruits & Vegetables	Onion	2500	16	3	2100
1	Fruits & Vegetables	Tomato Hybrid	4200	16	3	3500
2	Fruits & Vegetables	Tender Coconut	5100	15	3	4300
3	Fruits & Vegetables	Coriander Leaves	2000	15	3	1700
4	Fruits & Vegetables	Ladies Finger	1400	14	3	1200

In [18]: `Out_of_stock_items=df[df["outOfStock"]==True]
Out_of_stock_items[["name","quantity"]].sort_values(["quantity"],ascendin`

Out[18]:

	name	quantity
606	Everest Saffron Kesar	0
1120	Everest Saffron Kesar	0
1110	24 Mantra Organic Whole Wheat Atta	1
3625	Gala No Dust Broom	1
3623	Chakaachak Super Steel Scrubber	1
...
3076	Dettol Original Liquid Handwash Refill	750
3629	Comfort After Wash Fabric Conditioner Lily Fresh	860
2141	Del Monte Original Blend Tomato Ketchup Pouch	950
1753	Del Monte Original Blend Tomato Ketchup Pouch	950
2529	Del Monte Original Blend Tomato Ketchup Pouch	950


453 rows × 2 columns

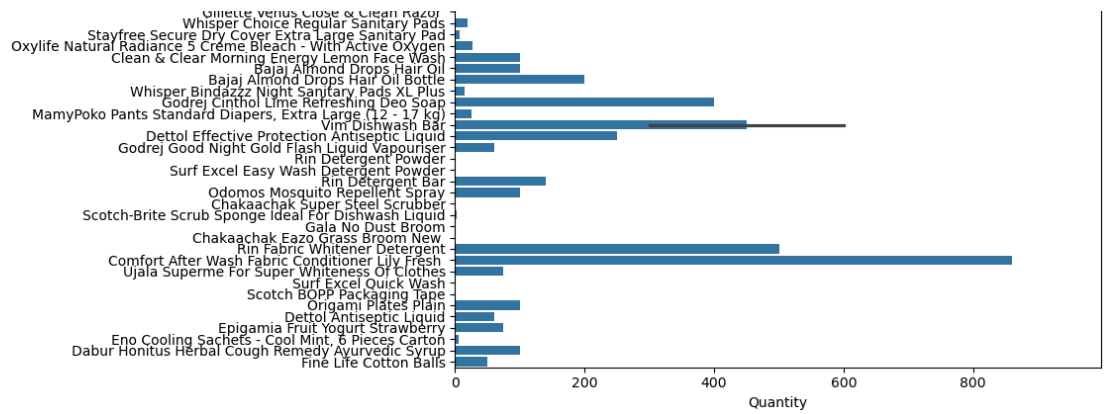
```
In [21]: plt.figure(figsize=(12, 25)) # Tall figure
sns.barplot(data=Out_of_stock_items, y="name", x="quantity", dodge=False)

plt.title("Out-of-Stock Items and Their Quantities")
plt.xlabel("Quantity")
plt.ylabel("Item Name")
plt.tight_layout()
plt.show()
```

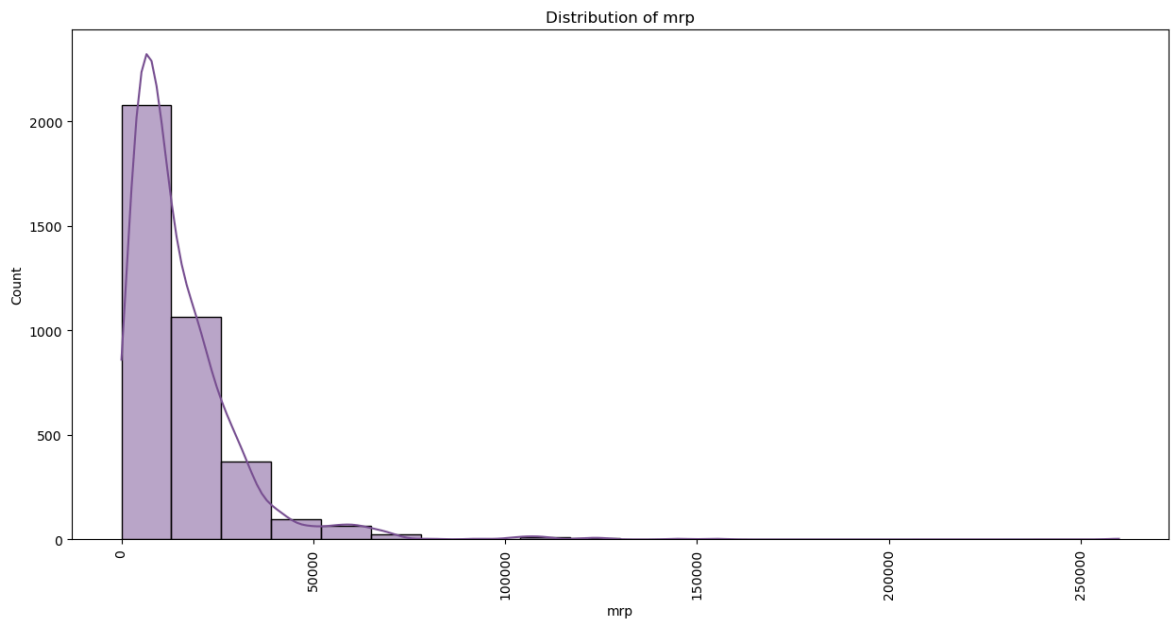
```
C:\Users\ashif\AppData\Local\Temp\ipykernel_13332\590317956.py:7: UserWarning: Glyph 150 (\x96) missing from font(s) DejaVu Sans.
  plt.tight_layout()
C:\Users\ashif\anaconda3\Lib\site-packages\IPython\core\pylabtools.py:170: UserWarning: Glyph 150 (\x96) missing from font(s) DejaVu Sans.
  fig.canvas.print_figure(bytes_io, **kw)
```

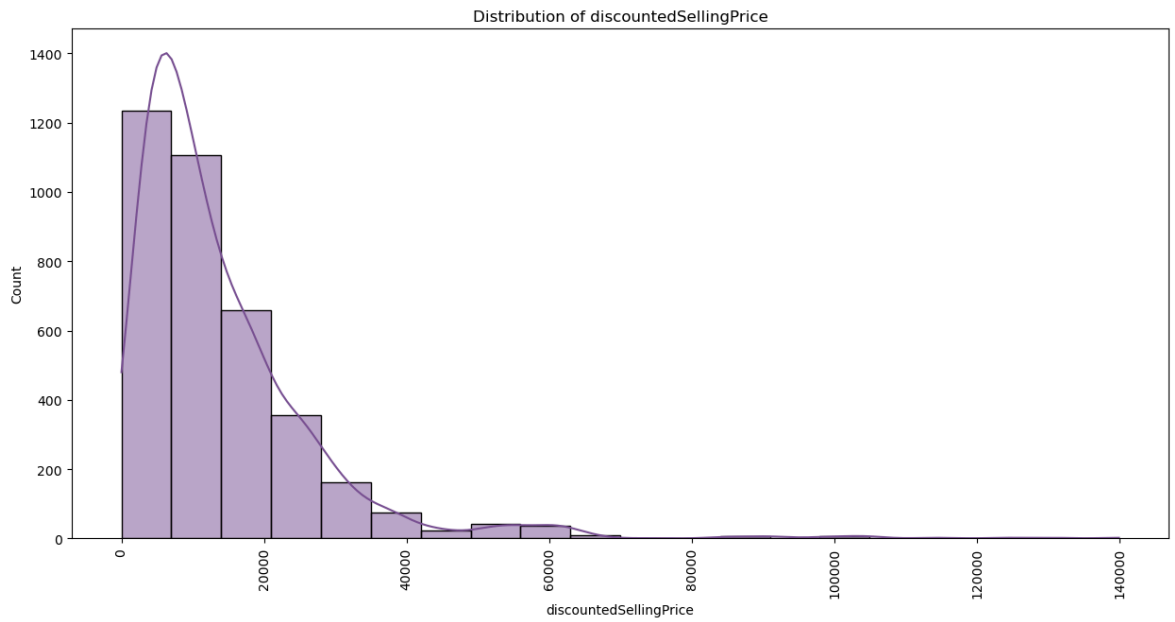
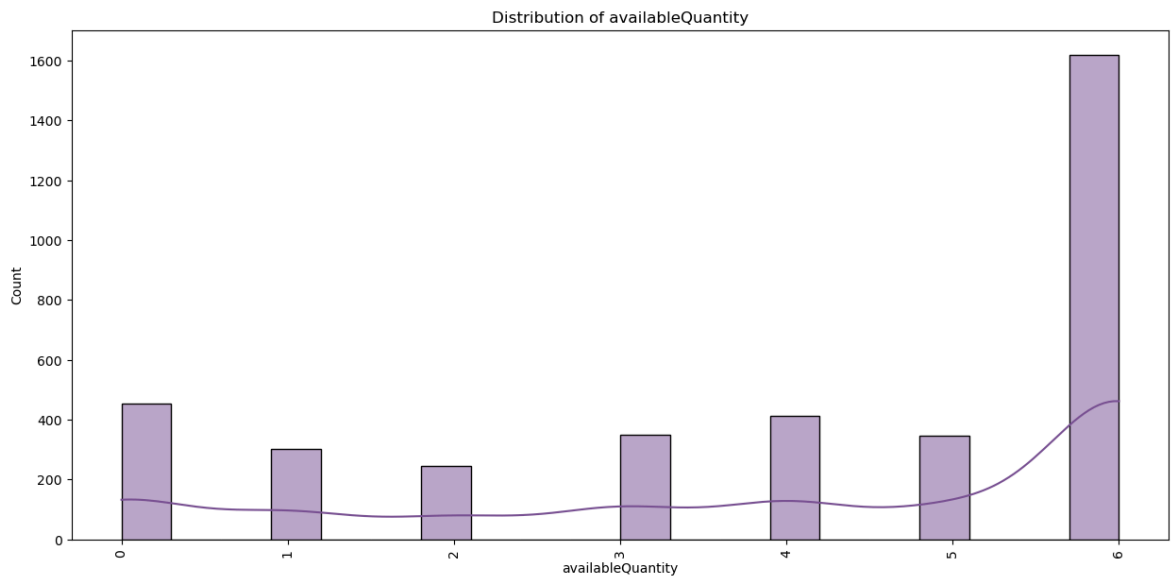
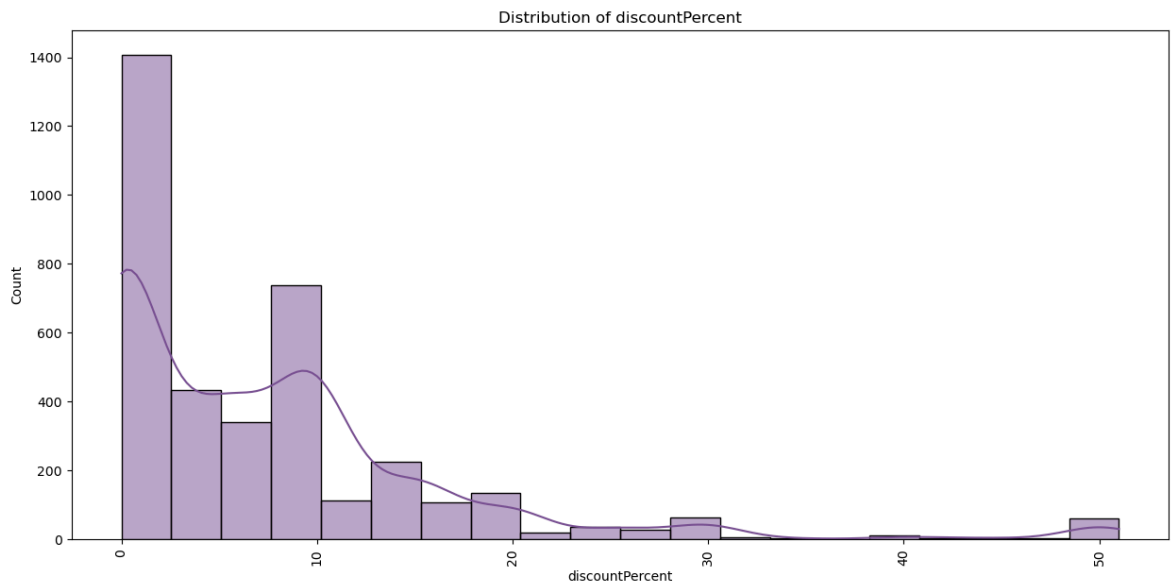
Out-of-Stock Items and Their Quantities


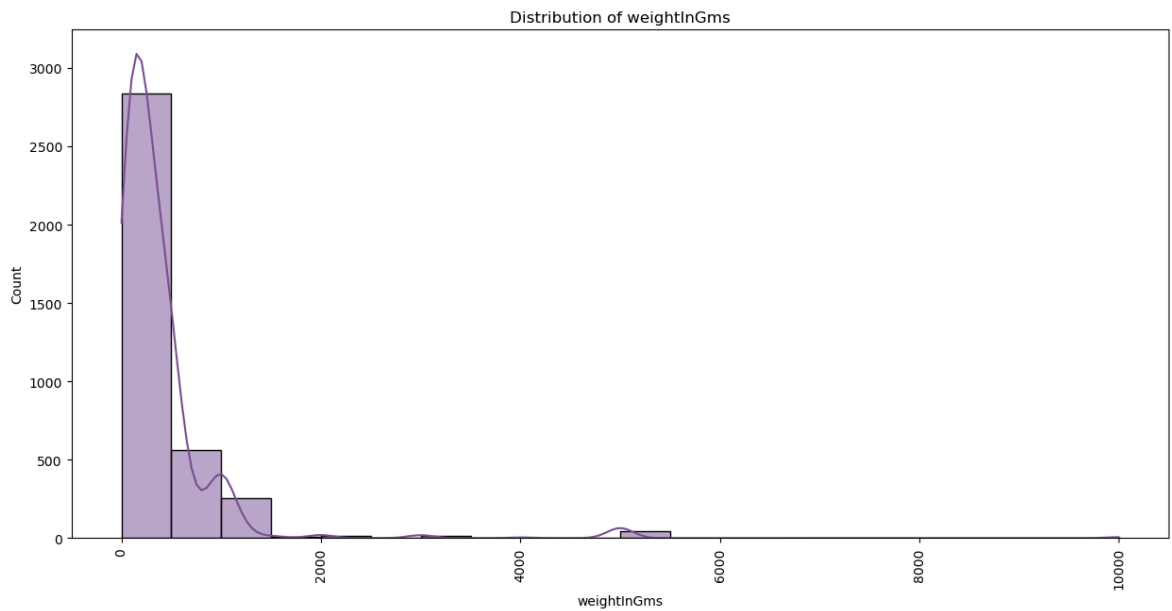




```
In [27]: d=df.select_dtypes(include='number')
for col in d:
    plt.figure(figsize=(15,7))
    sns.histplot(data=d,
x=col,
kde=True,
bins=20,
color='#7a5195')
    plt.title(f'Distribution of {col}')
    plt.xticks(rotation=90)
    plt.show()
```







In [31]: `df.head()`

Out[31]:

	Category	name	mrp	discountPercent	availableQuantity	discountedSellingPrice
0	Fruits & Vegetables	Onion	2500	16	3	2100
1	Fruits & Vegetables	Tomato Hybrid	4200	16	3	3500
2	Fruits & Vegetables	Tender Coconut	5100	15	3	4300
3	Fruits & Vegetables	Coriander Leaves	2000	15	3	1700
4	Fruits & Vegetables	Ladies Finger	1400	14	3	1200

In [33]: `Top_discount_percentage=df.sort_values(["discountPercent"],ascending=False)`
`Top_discount_percentage[["Category","name","mrp","discountPercent","discou`

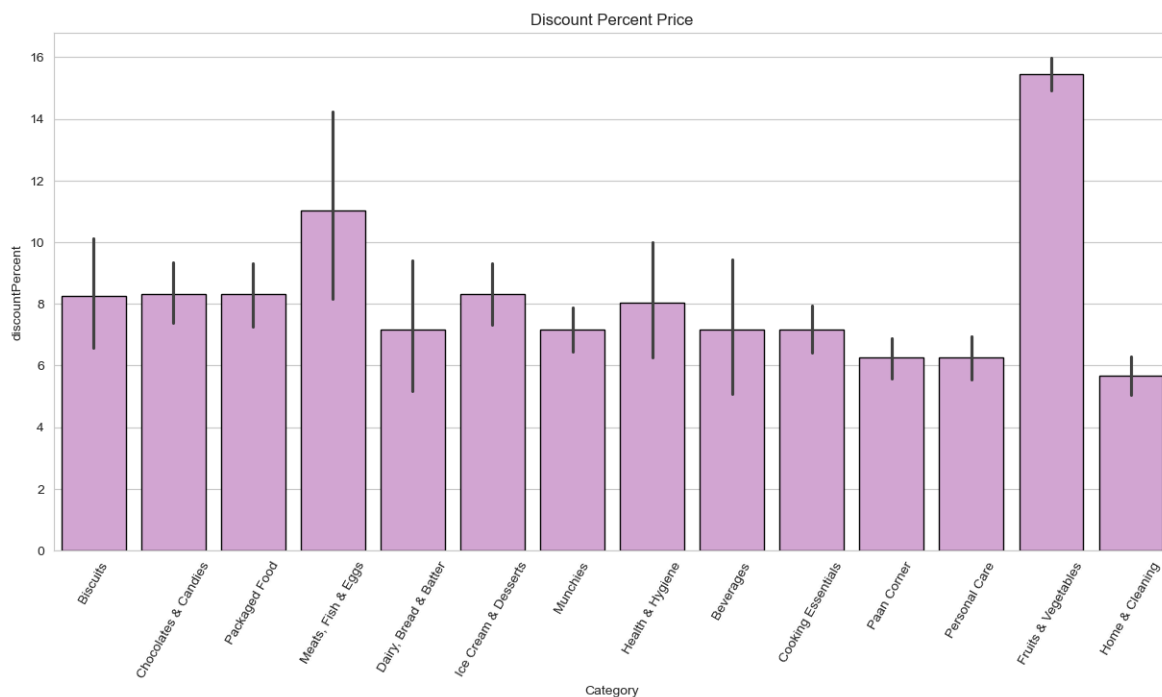
Out[33]:	Category	name	mrp	discountPercent	discountedSellingPrice
2608	Biscuits	Dukes Waffy Chocolate Wafers	4500	51	2200
2615	Biscuits	Dukes Waffy Orange Wafers	4500	51	2200
2619	Biscuits	Dukes Waffy Strawberry Wafers	4500	51	2200
2219	Chocolates & Candies	Chef's Basket Durum Wheat Fusilli Pasta	16000	50	8000
1431	Packaged Food	Chef's Basket Durum Wheat Elbow Pasta	16000	50	8000
...
2723	Biscuits	Britannia Marie Gold Biscuit	2800	0	2800
2722	Biscuits	Britannia Milk Bikis Milky Sandwich	2500	0	2500
1349	Beverages	Yoga Bar Peanut Butter Dark Chocolate Jar	24900	0	24900
2719	Biscuits	Britannia Milk Bikis Cream Biscuit	4500	0	4500
3731	Health & Hygiene	Dettol Antiseptic Liquid	3000	0	3000

3732 rows × 5 columns

```
In [39]: plt.figure(figsize=(15,7))
sns.set_style('whitegrid')
sns.barplot(data=Top_discount_percentage,
             x='Category',
             y='discountPercent',
             color='plum',
             edgecolor='black')

plt.xticks(rotation=60)
plt.title('Discount Percent Price')
```

Out[39]: Text(0.5, 1.0, 'Discount Percent Price')



Qunatity selling

```
In [40]: Top_quantity=df.sort_values(['quantity'])
Top_quantity[['Category','quantity','availableQuantity']]
```

```
Out[40]:
```

	Category	quantity	availableQuantity
3273	Paan Corner	0	3
2929	Personal Care	0	3
606	Cooking Essentials	0	0
3184	Paan Corner	0	2
2840	Personal Care	0	2
...
2320	Chocolates & Candies	1200	6
1544	Packaged Food	1200	6
1932	Ice Cream & Desserts	1200	6
3161	Paan Corner	1500	6
2817	Personal Care	1500	6


3732 rows × 3 columns

```
In [42]: plt.figure(figsize=(15,7))
sns.set_style('whitegrid')
sns.barplot(data=Top_quantity,
            x='Category',
            y='quantity',
            hue='availableQuantity',
            color='teal')
plt.xticks(rotation=60)
plt.title('Top quantity value');
```

C:\Users\ashif\AppData\Local\Temp\ipykernel_13332\1466617710.py:3: FutureWarning:


Setting a gradient palette using color= is deprecated and will be removed in v0.14.0. Set `palette='dark:teal'` for the same effect.

```
sns.barplot(data=Top_quantity,
```



Correlation Map

```
In [46]: plt.figure(figsize=(15,7))
corr=df.select_dtypes(include="number").corr()
sns.heatmap(data=corr,
            annot=True,
            fmt='.2f')
plt.show()
```



Applying in ANN model

Feature engineering

```
In [47]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3732 entries, 0 to 3731
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Category                             3732 non-null   object
1   name                                 3732 non-null   object
2   mrp                                  3732 non-null   int64
3   discountPercent                     3732 non-null   int64
4   availableQuantity                   3732 non-null   int64
5   discountedSellingPrice               3732 non-null   int64
6   weightInGms                         3732 non-null   int64
7   outOfStock                           3732 non-null   bool
8   quantity                             3732 non-null   int64
dtypes: bool(1), int64(6), object(2)
memory usage: 237.0+ KB
```

```
In [54]: label=['Category','name','outOfStock']
        from sklearn.preprocessing import LabelEncoder
```

```
In [55]: le=LabelEncoder()
```

```
In [56]: for feature in label:
        df[feature]=le.fit_transform(df[feature])
```

```
In [58]: df.head()
```

```
Out[58]:
```

	Category	name	mrp	discountPercent	availableQuantity	discountedSellingPrice	weightInGms
0	5	1046	2500	16	3	2100	
1	5	1551	4200	16	3	3500	
2	5	1500	5100	15	3	4300	
3	5	351	2000	15	3	1700	
4	5	835	1400	14	3	1200	

ANN Section (Artificial intelligence network)

```
In [60]: from sklearn.model_selection import train_test_split
        from sklearn.metrics import accuracy_score,recall_score,precision_score,f
```

```
In [61]: import tensorflow as tf
```

```
In [63]: ! pip install shap
```

```

Downloading shap-0.48.0-cp312-cp312-win_amd64.whl.metadata (25 kB)
Requirement already satisfied: numpy in c:\users\ashif\anaconda3\lib\site-packages (from shap) (1.26.4)
Requirement already satisfied: scipy in c:\users\ashif\anaconda3\lib\site-packages (from shap) (1.13.1)
Requirement already satisfied: scikit-learn in c:\users\ashif\anaconda3\lib\site-packages (from shap) (1.5.1)
Requirement already satisfied: pandas in c:\users\ashif\anaconda3\lib\site-packages (from shap) (2.2.3)
Requirement already satisfied: tqdm>=4.27.0 in c:\users\ashif\anaconda3\lib\site-packages (from shap) (4.66.5)
Requirement already satisfied: packaging>20.9 in c:\users\ashif\anaconda3\lib\site-packages (from shap) (24.1)
Collecting slicer==0.0.8 (from shap)
  Downloading slicer-0.0.8-py3-none-any.whl.metadata (4.0 kB)
Requirement already satisfied: numba>=0.54 in c:\users\ashif\anaconda3\lib\site-packages (from shap) (0.60.0)
Requirement already satisfied: cloudpickle in c:\users\ashif\anaconda3\lib\site-packages (from shap) (3.0.0)
Requirement already satisfied: typing-extensions in c:\users\ashif\anaconda3\lib\site-packages (from shap) (4.11.0)
Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in c:\users\ashif\anaconda3\lib\site-packages (from numba>=0.54->shap) (0.43.0)
Requirement already satisfied: colorama in c:\users\ashif\anaconda3\lib\site-packages (from tqdm>=4.27.0->shap) (0.4.6)
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\ashif\anaconda3\lib\site-packages (from pandas->shap) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in c:\users\ashif\anaconda3\lib\site-packages (from pandas->shap) (2024.1)
Requirement already satisfied: tzdata>=2022.7 in c:\users\ashif\anaconda3\lib\site-packages (from pandas->shap) (2023.3)
Requirement already satisfied: joblib>=1.2.0 in c:\users\ashif\anaconda3\lib\site-packages (from scikit-learn->shap) (1.4.2)
Requirement already satisfied: threadpoolctl>=3.1.0 in c:\users\ashif\anaconda3\lib\site-packages (from scikit-learn->shap) (3.5.0)
Requirement already satisfied: six>=1.5 in c:\users\ashif\anaconda3\lib\site-packages (from python-dateutil>=2.8.2->pandas->shap) (1.16.0)
Downloading shap-0.48.0-cp312-cp312-win_amd64.whl (545 kB)
----- 0.0/545.3 kB ? eta -:--:--
----- 545.3/545.3 kB 6.0 MB/s eta 0:00:00
Downloading slicer-0.0.8-py3-none-any.whl (15 kB)
Installing collected packages: slicer, shap
Successfully installed shap-0.48.0 slicer-0.0.8


```


[illegible]


```
In [82]: scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```


```
In [83]: model=Sequential([
    Input(shape=(X_train_scaled.shape[1],)),
    Dense(128,activation='relu'),
    Dense(64,activation='relu'),
    Dense(1)
])
```


```
In [86]: model.compile(optimizer=Adam(),
    loss='mean_squared_error',
    metrics=['mse'])
model.fit(X_train_scaled,
    y_train,epochs=100,
    batch_size=32,
    validation_split=0.1)
```



Epoch 1/100
84/84  1s 4ms/step - loss: 0.0046 - mse: 0.0046 - val_loss: 0.0030 - val_mse: 0.0030


Epoch 2/100
84/84  0s 2ms/step - loss: 0.0021 - mse: 0.0021 - val_loss: 0.0021 - val_mse: 0.0021


Epoch 3/100
84/84  0s 2ms/step - loss: 0.0017 - mse: 0.0017 - val_loss: 0.0015 - val_mse: 0.0015


Epoch 4/100
84/84  0s 2ms/step - loss: 0.0018 - mse: 0.0018 - val_loss: 0.0021 - val_mse: 0.0021


Epoch 5/100
84/84  0s 2ms/step - loss: 0.0021 - mse: 0.0021 - val_loss: 0.0012 - val_mse: 0.0012


Epoch 6/100
84/84  0s 2ms/step - loss: 0.0010 - mse: 0.0010 - val_loss: 8.3732e-04 - val_mse: 8.3732e-04


Epoch 7/100
84/84  0s 2ms/step - loss: 7.3045e-04 - mse: 7.3045e-04 - val_loss: 9.2200e-04 - val_mse: 9.2200e-04


Epoch 8/100
84/84  0s 2ms/step - loss: 9.1120e-04 - mse: 9.1120e-04 - val_loss: 0.0010 - val_mse: 0.0010


Epoch 9/100
84/84  0s 2ms/step - loss: 5.9457e-04 - mse: 5.9457e-04 - val_loss: 0.0010 - val_mse: 0.0010


Epoch 10/100
84/84  0s 2ms/step - loss: 6.3691e-04 - mse: 6.3691e-04 - val_loss: 7.6235e-04 - val_mse: 7.6235e-04


Epoch 11/100
84/84  0s 2ms/step - loss: 4.9928e-04 - mse: 4.9928e-04 - val_loss: 6.3912e-04 - val_mse: 6.3912e-04


Epoch 12/100
84/84  0s 2ms/step - loss: 3.9446e-04 - mse: 3.9446e-04 - val_loss: 5.4683e-04 - val_mse: 5.4683e-04


Epoch 13/100
84/84  0s 3ms/step - loss: 4.9065e-04 - mse: 4.9065e-04 - val_loss: 7.1757e-04 - val_mse: 7.1757e-04


Epoch 14/100
84/84  0s 2ms/step - loss: 5.3326e-04 - mse: 5.3326e-04 - val_loss: 5.2001e-04 - val_mse: 5.2001e-04


Epoch 15/100
84/84  0s 2ms/step - loss: 2.8861e-04 - mse: 2.8861e-04 - val_loss: 7.3625e-04 - val_mse: 7.3625e-04

Epoch 16/100
84/84  0s 2ms/step - loss: 0.0018 - mse: 0.0018 - val_loss: 0.0015 - val_mse: 0.0015

Epoch 17/100
84/84  0s 2ms/step - loss: 8.8839e-04 - mse: 8.8839e-04 - val_loss: 4.6048e-04 - val_mse: 4.6048e-04

Epoch 18/100
84/84  0s 2ms/step - loss: 4.9846e-04 - mse: 4.9846e-04 - val_loss: 8.3501e-04 - val_mse: 8.3501e-04

Epoch 19/100
84/84  0s 2ms/step - loss: 4.7885e-04 - mse: 4.7885e-04 - val_loss: 6.3975e-04 - val_mse: 6.3975e-04

Epoch 20/100
84/84  0s 2ms/step - loss: 3.3272e-04 - mse: 3.3272e-04 - val_loss: 3.8675e-04 - val_mse: 3.8675e-04

Epoch 21/100
84/84 ————— 0s 2ms/step - loss: 2.9285e-04 - mse: 2.9285e-04 - val_loss: 3.3543e-04 - val_mse: 3.3543e-04
Epoch 22/100
84/84 ————— 0s 2ms/step - loss: 2.5958e-04 - mse: 2.5958e-04 - val_loss: 5.6141e-04 - val_mse: 5.6141e-04
Epoch 23/100
84/84 ————— 0s 2ms/step - loss: 4.2862e-04 - mse: 4.2862e-04 - val_loss: 5.0637e-04 - val_mse: 5.0637e-04
Epoch 24/100
84/84 ————— 0s 3ms/step - loss: 2.8760e-04 - mse: 2.8760e-04 - val_loss: 0.0016 - val_mse: 0.0016
Epoch 25/100
84/84 ————— 0s 2ms/step - loss: 8.8997e-04 - mse: 8.8997e-04 - val_loss: 5.4797e-04 - val_mse: 5.4797e-04
Epoch 26/100
84/84 ————— 0s 2ms/step - loss: 5.5122e-04 - mse: 5.5122e-04 - val_loss: 4.2937e-04 - val_mse: 4.2937e-04
Epoch 27/100
84/84 ————— 0s 2ms/step - loss: 3.7826e-04 - mse: 3.7826e-04 - val_loss: 8.7348e-04 - val_mse: 8.7348e-04
Epoch 28/100
84/84 ————— 0s 2ms/step - loss: 4.6969e-04 - mse: 4.6969e-04 - val_loss: 8.8197e-04 - val_mse: 8.8197e-04
Epoch 29/100
84/84 ————— 0s 2ms/step - loss: 7.3658e-04 - mse: 7.3658e-04 - val_loss: 5.7461e-04 - val_mse: 5.7461e-04
Epoch 30/100
84/84 ————— 0s 2ms/step - loss: 5.7704e-04 - mse: 5.7704e-04 - val_loss: 6.5893e-04 - val_mse: 6.5893e-04
Epoch 31/100
84/84 ————— 0s 2ms/step - loss: 4.9467e-04 - mse: 4.9467e-04 - val_loss: 5.4343e-04 - val_mse: 5.4343e-04
Epoch 32/100
84/84 ————— 0s 2ms/step - loss: 3.2725e-04 - mse: 3.2725e-04 - val_loss: 3.8524e-04 - val_mse: 3.8524e-04
Epoch 33/100
84/84 ————— 0s 2ms/step - loss: 1.8227e-04 - mse: 1.8227e-04 - val_loss: 3.0552e-04 - val_mse: 3.0552e-04
Epoch 34/100
84/84 ————— 0s 2ms/step - loss: 2.0664e-04 - mse: 2.0664e-04 - val_loss: 2.7737e-04 - val_mse: 2.7737e-04
Epoch 35/100
84/84 ————— 0s 2ms/step - loss: 1.8089e-04 - mse: 1.8089e-04 - val_loss: 1.6250e-04 - val_mse: 1.6250e-04
Epoch 36/100
84/84 ————— 0s 2ms/step - loss: 1.0679e-04 - mse: 1.0679e-04 - val_loss: 3.6486e-04 - val_mse: 3.6486e-04
Epoch 37/100
84/84 ————— 0s 2ms/step - loss: 1.5951e-04 - mse: 1.5951e-04 - val_loss: 9.6381e-04 - val_mse: 9.6381e-04
Epoch 38/100
84/84 ————— 0s 2ms/step - loss: 6.5370e-04 - mse: 6.5370e-04 - val_loss: 3.0382e-04 - val_mse: 3.0382e-04
Epoch 39/100
84/84 ————— 0s 2ms/step - loss: 7.5670e-04 - mse: 7.5670e-04 - val_loss: 2.5439e-04 - val_mse: 2.5439e-04
Epoch 40/100
84/84 ————— 0s 3ms/step - loss: 1.2392e-04 - mse: 1.2392e-04 - val_loss: 3.7029e-04 - val_mse: 3.7029e-04

Epoch 41/100
84/84 ————— 0s 2ms/step - loss: 1.7307e-04 - mse: 1.7307e-04 - val_loss: 3.4336e-04 - val_mse: 3.4336e-04
Epoch 42/100
84/84 ————— 0s 2ms/step - loss: 3.2646e-04 - mse: 3.2646e-04 - val_loss: 2.2832e-04 - val_mse: 2.2832e-04
Epoch 43/100
84/84 ————— 0s 2ms/step - loss: 2.0150e-04 - mse: 2.0150e-04 - val_loss: 3.2426e-04 - val_mse: 3.2426e-04
Epoch 44/100
84/84 ————— 0s 2ms/step - loss: 2.7754e-04 - mse: 2.7754e-04 - val_loss: 1.6434e-04 - val_mse: 1.6434e-04
Epoch 45/100
84/84 ————— 0s 2ms/step - loss: 1.9279e-04 - mse: 1.9279e-04 - val_loss: 6.4458e-04 - val_mse: 6.4458e-04
Epoch 46/100
84/84 ————— 0s 2ms/step - loss: 4.5137e-04 - mse: 4.5137e-04 - val_loss: 6.6478e-04 - val_mse: 6.6478e-04
Epoch 47/100
84/84 ————— 0s 2ms/step - loss: 2.9063e-04 - mse: 2.9063e-04 - val_loss: 6.4317e-04 - val_mse: 6.4317e-04
Epoch 48/100
84/84 ————— 0s 2ms/step - loss: 5.4536e-04 - mse: 5.4536e-04 - val_loss: 9.8531e-04 - val_mse: 9.8531e-04
Epoch 49/100
84/84 ————— 0s 2ms/step - loss: 0.0011 - mse: 0.0011 - val_loss: 2.5147e-04 - val_mse: 2.5147e-04
Epoch 50/100
84/84 ————— 0s 2ms/step - loss: 1.8629e-04 - mse: 1.8629e-04 - val_loss: 7.5062e-04 - val_mse: 7.5062e-04
Epoch 51/100
84/84 ————— 0s 2ms/step - loss: 3.5871e-04 - mse: 3.5871e-04 - val_loss: 1.3777e-04 - val_mse: 1.3777e-04
Epoch 52/100
84/84 ————— 0s 2ms/step - loss: 1.7069e-04 - mse: 1.7069e-04 - val_loss: 2.0620e-04 - val_mse: 2.0620e-04
Epoch 53/100
84/84 ————— 0s 2ms/step - loss: 1.1897e-04 - mse: 1.1897e-04 - val_loss: 1.8816e-04 - val_mse: 1.8816e-04
Epoch 54/100
84/84 ————— 0s 2ms/step - loss: 1.8092e-04 - mse: 1.8092e-04 - val_loss: 2.2791e-04 - val_mse: 2.2791e-04
Epoch 55/100
84/84 ————— 0s 2ms/step - loss: 2.9251e-04 - mse: 2.9251e-04 - val_loss: 2.3388e-04 - val_mse: 2.3388e-04
Epoch 56/100
84/84 ————— 0s 2ms/step - loss: 2.2887e-04 - mse: 2.2887e-04 - val_loss: 7.5892e-04 - val_mse: 7.5892e-04
Epoch 57/100
84/84 ————— 0s 2ms/step - loss: 4.8976e-04 - mse: 4.8976e-04 - val_loss: 1.8976e-04 - val_mse: 1.8976e-04
Epoch 58/100
84/84 ————— 0s 2ms/step - loss: 1.1326e-04 - mse: 1.1326e-04 - val_loss: 1.9471e-04 - val_mse: 1.9471e-04
Epoch 59/100
84/84 ————— 0s 2ms/step - loss: 1.1549e-04 - mse: 1.1549e-04 - val_loss: 3.2531e-04 - val_mse: 3.2531e-04
Epoch 60/100
84/84 ————— 0s 2ms/step - loss: 9.4978e-05 - mse: 9.4978e-05 - val_loss: 2.7650e-04 - val_mse: 2.7650e-04

Epoch 61/100
84/84 ————— 0s 2ms/step - loss: 2.4203e-04 - mse: 2.4203e-04 - val_loss: 2.2948e-04 - val_mse: 2.2948e-04
Epoch 62/100
84/84 ————— 0s 2ms/step - loss: 1.4826e-04 - mse: 1.4826e-04 - val_loss: 4.8371e-04 - val_mse: 4.8371e-04
Epoch 63/100
84/84 ————— 0s 2ms/step - loss: 2.7365e-04 - mse: 2.7365e-04 - val_loss: 1.4029e-04 - val_mse: 1.4029e-04
Epoch 64/100
84/84 ————— 0s 2ms/step - loss: 7.7230e-05 - mse: 7.7230e-05 - val_loss: 1.8268e-04 - val_mse: 1.8268e-04
Epoch 65/100
84/84 ————— 0s 2ms/step - loss: 6.2853e-05 - mse: 6.2853e-05 - val_loss: 2.0831e-04 - val_mse: 2.0831e-04
Epoch 66/100
84/84 ————— 0s 2ms/step - loss: 2.3778e-04 - mse: 2.3778e-04 - val_loss: 1.5975e-04 - val_mse: 1.5975e-04
Epoch 67/100
84/84 ————— 0s 2ms/step - loss: 1.1240e-04 - mse: 1.1240e-04 - val_loss: 1.3893e-04 - val_mse: 1.3893e-04
Epoch 68/100
84/84 ————— 0s 2ms/step - loss: 8.3955e-05 - mse: 8.3955e-05 - val_loss: 2.5568e-04 - val_mse: 2.5568e-04
Epoch 69/100
84/84 ————— 0s 2ms/step - loss: 1.0388e-04 - mse: 1.0388e-04 - val_loss: 3.3471e-04 - val_mse: 3.3471e-04
Epoch 70/100
84/84 ————— 0s 2ms/step - loss: 3.5332e-04 - mse: 3.5332e-04 - val_loss: 2.6968e-04 - val_mse: 2.6968e-04
Epoch 71/100
84/84 ————— 0s 2ms/step - loss: 2.6177e-04 - mse: 2.6177e-04 - val_loss: 1.6623e-04 - val_mse: 1.6623e-04
Epoch 72/100
84/84 ————— 0s 2ms/step - loss: 1.1748e-04 - mse: 1.1748e-04 - val_loss: 2.1731e-04 - val_mse: 2.1731e-04
Epoch 73/100
84/84 ————— 0s 3ms/step - loss: 1.8857e-04 - mse: 1.8857e-04 - val_loss: 6.7454e-04 - val_mse: 6.7454e-04
Epoch 74/100
84/84 ————— 0s 3ms/step - loss: 2.6078e-04 - mse: 2.6078e-04 - val_loss: 2.5877e-04 - val_mse: 2.5877e-04
Epoch 75/100
84/84 ————— 0s 3ms/step - loss: 4.2363e-04 - mse: 4.2363e-04 - val_loss: 3.9799e-04 - val_mse: 3.9799e-04
Epoch 76/100
84/84 ————— 0s 3ms/step - loss: 4.5179e-04 - mse: 4.5179e-04 - val_loss: 5.5317e-04 - val_mse: 5.5317e-04
Epoch 77/100
84/84 ————— 0s 4ms/step - loss: 3.3746e-04 - mse: 3.3746e-04 - val_loss: 4.2767e-04 - val_mse: 4.2767e-04
Epoch 78/100
84/84 ————— 0s 3ms/step - loss: 1.7937e-04 - mse: 1.7937e-04 - val_loss: 2.4753e-04 - val_mse: 2.4753e-04
Epoch 79/100
84/84 ————— 0s 3ms/step - loss: 2.6018e-04 - mse: 2.6018e-04 - val_loss: 1.8837e-04 - val_mse: 1.8837e-04
Epoch 80/100
84/84 ————— 0s 3ms/step - loss: 9.4656e-05 - mse: 9.4656e-05 - val_loss: 6.4708e-05 - val_mse: 6.4708e-05

Epoch 81/100
84/84 ————— 0s 3ms/step - loss: 6.5077e-05 - mse: 6.5077e-05 - val_loss: 4.8466e-05 - val_mse: 4.8466e-05
Epoch 82/100
84/84 ————— 0s 4ms/step - loss: 4.5573e-05 - mse: 4.5573e-05 - val_loss: 1.2191e-04 - val_mse: 1.2191e-04
Epoch 83/100
84/84 ————— 0s 3ms/step - loss: 6.7483e-05 - mse: 6.7483e-05 - val_loss: 4.5132e-05 - val_mse: 4.5132e-05
Epoch 84/100
84/84 ————— 0s 3ms/step - loss: 4.3571e-05 - mse: 4.3571e-05 - val_loss: 9.3072e-05 - val_mse: 9.3072e-05
Epoch 85/100
84/84 ————— 0s 2ms/step - loss: 6.6859e-05 - mse: 6.6859e-05 - val_loss: 1.4494e-04 - val_mse: 1.4494e-04
Epoch 86/100
84/84 ————— 0s 2ms/step - loss: 9.7507e-05 - mse: 9.7507e-05 - val_loss: 1.3516e-04 - val_mse: 1.3516e-04
Epoch 87/100
84/84 ————— 0s 2ms/step - loss: 1.3500e-04 - mse: 1.3500e-04 - val_loss: 5.9554e-05 - val_mse: 5.9554e-05
Epoch 88/100
84/84 ————— 0s 2ms/step - loss: 8.1853e-05 - mse: 8.1853e-05 - val_loss: 1.1984e-04 - val_mse: 1.1984e-04
Epoch 89/100
84/84 ————— 0s 3ms/step - loss: 9.4043e-05 - mse: 9.4043e-05 - val_loss: 1.3820e-04 - val_mse: 1.3820e-04
Epoch 90/100
84/84 ————— 0s 3ms/step - loss: 1.2542e-04 - mse: 1.2542e-04 - val_loss: 1.9397e-04 - val_mse: 1.9397e-04
Epoch 91/100
84/84 ————— 0s 4ms/step - loss: 1.9652e-04 - mse: 1.9652e-04 - val_loss: 1.8354e-04 - val_mse: 1.8354e-04
Epoch 92/100
84/84 ————— 0s 2ms/step - loss: 8.2534e-05 - mse: 8.2534e-05 - val_loss: 2.3352e-04 - val_mse: 2.3352e-04
Epoch 93/100
84/84 ————— 0s 2ms/step - loss: 2.2416e-04 - mse: 2.2416e-04 - val_loss: 0.0019 - val_mse: 0.0019
Epoch 94/100
84/84 ————— 0s 3ms/step - loss: 0.0013 - mse: 0.0013 - val_loss: 7.2335e-04 - val_mse: 7.2335e-04
Epoch 95/100
84/84 ————— 0s 2ms/step - loss: 7.0522e-04 - mse: 7.0522e-04 - val_loss: 4.7542e-04 - val_mse: 4.7542e-04
Epoch 96/100
84/84 ————— 0s 2ms/step - loss: 2.7387e-04 - mse: 2.7387e-04 - val_loss: 2.1938e-04 - val_mse: 2.1938e-04
Epoch 97/100
84/84 ————— 0s 2ms/step - loss: 1.1216e-04 - mse: 1.1216e-04 - val_loss: 3.8655e-04 - val_mse: 3.8655e-04
Epoch 98/100
84/84 ————— 0s 2ms/step - loss: 2.0430e-04 - mse: 2.0430e-04 - val_loss: 1.5194e-04 - val_mse: 1.5194e-04
Epoch 99/100
84/84 ————— 0s 2ms/step - loss: 9.5175e-05 - mse: 9.5175e-05 - val_loss: 1.6619e-04 - val_mse: 1.6619e-04
Epoch 100/100
84/84 ————— 0s 2ms/step - loss: 1.0510e-04 - mse: 1.0510e-04 - val_loss: 8.1086e-05 - val_mse: 8.1086e-05

Out[86]: <keras.src.callbacks.history.History at 0x1db75a78bf0>

```
In [87]: y_pred = model.predict(X_test_scaled)
y_pred_labels = (y_pred > 0.5).astype(int).flatten()
```

24/24 ————— 0s 3ms/step

```
In [88]: print(classification_report(y_test,y_pred_labels))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	648
1	1.00	1.00	1.00	99
accuracy			1.00	747
macro avg	1.00	1.00	1.00	747
weighted avg	1.00	1.00	1.00	747

```
In [90]: # Convert probabilities to class labels
y_train_pred = model.predict(X_train)
y_test_pred = model.predict(X_test)

# Convert probabilities to class labels
y_train_pred_labels = (y_train_pred > 0.5).astype(int)
y_test_pred_labels = (y_test_pred > 0.5).astype(int)

# Evaluate
from sklearn.metrics import classification_report

print("Train Report:")
print(classification_report(y_train, y_train_pred_labels))

print("Test Report:")
print(classification_report(y_test, y_test_pred_labels))
```

94/94 ————— 0s 1ms/step

24/24 ————— 0s 1ms/step

Train Report:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	2631
1	0.12	1.00	0.21	354
accuracy			0.12	2985
macro avg	0.06	0.50	0.11	2985
weighted avg	0.01	0.12	0.03	2985

Test Report:

	precision	recall	f1-score	support
0	0.00	0.00	0.00	648
1	0.13	1.00	0.23	99
accuracy			0.13	747
macro avg	0.07	0.50	0.12	747
weighted avg	0.02	0.13	0.03	747

```

C:\Users\ashif\anaconda3\Lib\site-packages\sklearn\metrics\_classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
C:\Users\ashif\anaconda3\Lib\site-packages\sklearn\metrics\_classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
C:\Users\ashif\anaconda3\Lib\site-packages\sklearn\metrics\_classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
C:\Users\ashif\anaconda3\Lib\site-packages\sklearn\metrics\_classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
C:\Users\ashif\anaconda3\Lib\site-packages\sklearn\metrics\_classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
C:\Users\ashif\anaconda3\Lib\site-packages\sklearn\metrics\_classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

Conclusion:

- Current ANN is severely underperforming and overfitting to the minority class.
- You should address class imbalance and retrain with class_weight, SMOTE, or both.
- After fixing, re-run the classification report to see improvements in both recall and precision across both classes.

In []: